

Supplementary Material: Plato-Bio

A reproducibility inventory for the structural-biology case study and deterministic software validation.

S1. Reproduction commands

- `.venv/bin/python preprint/experiments/run_globin_structure_benchmark.py`
- `.venv/bin/python preprint/experiments/run_software_validation.py`
- `.venv/bin/python preprint/experiments/build_summary_figures.py`

S2. Structural input inventory

File	Bytes	SHA-256
1A3N.pdb	461619	1ce17adb5a0a4f986e173aaffa3c3e9910f483f9e32361615a75abb325a1aa7f
3RGK.pdb	147015	5fdbab6e63ee6d857f3d31cf639452ade4db527c6b14289a13b6a3aa11660f2a
AF-P02144-F1.pdb	102950	68dd0c9976dd37bc6787b5b398ddc3ac03b6abdb7cc7025298dd503d8bb3c882
AF-P68871-F1.pdb	96551	37b411f5eaa2ff2e6f51774da344f1bd7b54e78c98ce67d9ffb02d8a3dbdf32a
AF-P69905-F1.pdb	92096	d9d7ad4607164de17b65767b467d366909d74f96df5176ee368a9a515aaa1911

S3. Target-level structural results

Target	Matched	Identity	RMSD Å	Median Å	≤2 Å	p	P
hemoglobin_alpha	141	1.000	0.270	0.220	1.000	0.242	0.00388
hemoglobin_beta	145	1.000	0.520	0.453	1.000	0.353	1.34e-05
myoglobin	147	0.987	0.501	0.310	0.993	0.283	0.000515

The residue-level file contains 433 matched positions. It is supplied as `residue_metrics.csv` rather than reproduced as a 433-row PDF table.

S4. Deterministic validation results

Suite	Tests	Passed	Skipped	Failures	Errors	Wall s
adversarial_safety	57	56	1	0	0	0.89
biology_domain	12	12	0	0	0	0.70

Suite	Tests	Passed	Skipped	Failures	Errors	Wall s
evidence_and_citations	46	46	0	0	0	1.54
full_python	918	912	6	0	0	67.63
genomics_adapters	11	11	0	0	0	1.15

Targeted suites overlap with the full Python suite. Counts are not additive. Skips remain unvalidated live or platform-specific paths.

S5. Measurement-repair acceptance criteria

- A GoldenTask with domain=biology constructs Plato with domain=biology.
- Method-signal recall is computed from methods.md and appears in summary metrics.
- Evidence JSONL persists every drafted Claim row before emitted EvidenceLink rows.
- A drafted claim with no supporting source remains in the denominator and yields unsupported_claim_rate=1.0.
- A synthetic rigid rotation and translation is recovered by the Kabsch implementation within numerical tolerance.

S6. Interpretation and unavailable live lanes

The default evaluation runner executes idea and method stages only. It does not execute results or paper generation in this release. Live LLM, E2B, Modal, hosted PostgreSQL, and authenticated Hugging Face evaluations require external credentials or services and were not inferred from deterministic tests. Same-model reviewer roles are self-critique, not peer review. The autonomous-loop adapters do not establish autonomous scientific improvement.

S7. bioRxiv packaging notes

- Main manuscript is supplied as one PDF with embedded figures and tables.
- Supplemental data are separate from the main manuscript.
- The work is an English-language computational methods/validation article with Results and Discussion.
- No clinical, human-subject, animal, or identifiable-person data are included.
- Generative AI assistance is disclosed; the human author remains accountable and the AI is not an author.
- The related Denario arXiv preprint is cited as prior work and is not duplicated by this study.